

fmincon

Find a minimum of a constrained nonlinear multivariable function

$$\min_x f(x) \quad \text{subject to}$$

$$c(x) \leq 0$$

$$ceq(x) = 0$$

$$A \cdot x \leq b$$

$$Aeq \cdot x = beq$$

$$lb \leq x \leq ub$$

where x , b , beq , lb , and ub are vectors, A and Aeq are matrices, $c(x)$ and $ceq(x)$ are functions that return vectors, and $f(x)$ is a function that returns a scalar. $f(x)$, $c(x)$, and $ceq(x)$ can be nonlinear functions.

Syntax

```
x = fmincon(fun,x0,A,b)
x = fmincon(fun,x0,A,b,Aeq,beq)
x = fmincon(fun,x0,A,b,Aeq,beq,lb,ub)
x = fmincon(fun,x0,A,b,Aeq,beq,lb,ub,nonlcon)
x = fmincon(fun,x0,A,b,Aeq,beq,lb,ub,nonlcon,options)
x = fmincon(fun,x0,A,b,Aeq,beq,lb,ub,nonlcon,options,P1,P2,...)
[x,fval] = fmincon(...)
[x,fval,exitflag] = fmincon(...)
[x,fval,exitflag,output] = fmincon(...)
[x,fval,exitflag,output,lambda] = fmincon(...)
[x,fval,exitflag,output,lambda,grad] = fmincon(...)
[x,fval,exitflag,output,lambda,grad,hessian] = fmincon(...)
```

Description

fmincon finds a constrained minimum of a scalar function of several variables starting at an initial estimate. This is generally referred to as *constrained nonlinear optimization* or *nonlinear programming*.

`x = fmincon(fun,x0,A,b)` starts at `x0` and finds a minimum `x` to the function described in `fun` subject to the linear inequalities $A \cdot x \leq b$. `x0` can be a scalar, vector, or matrix.

`x = fmincon(fun,x0,A,b,Aeq,beq)` minimizes `fun` subject to the linear equalities $Aeq \cdot x = beq$ as well as $A \cdot x \leq b$. Set `A=[]` and `b=[]` if no inequalities exist.

`x = fmincon(fun,x0,A,b,Aeq,beq,lb,ub)` defines a set of lower and upper bounds on the design variables, `x`, so that the solution is always in the range $lb \leq x \leq ub$. Set `Aeq=[]` and `beq=[]` if no equalities exist.

`x = fmincon(fun,x0,A,b,Aeq,beq,lb,ub,nonlcon)` subjects the minimization to the nonlinear inequalities `c(x)` or equalities `ceq(x)` defined in `nonlcon`. `fmincon` optimizes such that $c(x) \leq 0$ and $ceq(x) = 0$. Set `lb=[]` and/or `ub=[]` if no bounds exist.

`x = fmincon(fun,x0,A,b,Aeq,beq,lb,ub,nonlcon,options)` minimizes with the optimization parameters specified in the structure `options`. Use `optimset` to set these parameters.

`x = fmincon(fun,x0,A,b,Aeq,beq,lb,ub,nonlcon,options,P1,P2,...)` passes the problem-dependent parameters `P1`, `P2`, etc., directly to the functions `fun` and `nonlcon`. Pass empty matrices as placeholders for `A`, `b`, `Aeq`, `beq`, `lb`, `ub`, `nonlcon`, and `options` if these arguments are not needed.

`[x,fval] = fmincon(...)` returns the value of the objective function `fun` at the solution `x`.

`[x,fval,exitflag] = fmincon(...)` returns a value `exitflag` that describes the exit condition of `fmincon`.

`[x,fval,exitflag,output] = fmincon(...)` returns a structure `output` with information about the optimization.

`[x,fval,exitflag,output,lambda] = fmincon(...)` returns a structure `lambda` whose fields contain the Lagrange multipliers at the solution `x`.

`[x,fval,exitflag,output,lambda,grad] = fmincon(...)` returns the value of the gradient of `fun` at the solution `x`.

`[x,fval,exitflag,output,lambda,grad,hessian] = fmincon(...)` returns the value of the Hessian of `fun` at the solution `x`.

Input Arguments

[Function Arguments](#) contains general descriptions of arguments passed in to `fmincon`. This "Arguments" section provides function-specific details for `fun`, `nonlcon`, and `options`:

`fun` The function to be minimized. `fun` is a function that accepts a scalar `x` and returns a scalar `f`, the objective function evaluated at `x`. The function `fun` can be specified as a function handle.

```
x = fmincon(@myfun,x0,A,b)
```

where `myfun` is a MATLAB function such as

```
function f = myfun(x)
f = ...           % Compute function value at x
```

`fun` can also be an inline object.

```
x = fmincon(inline('norm(x)^2'),x0,A,b);
```

If the gradient of `fun` can also be computed *and* the `GradObj` parameter is 'on', as set by

```
options = optimset('GradObj','on')
```

then the function `fun` must return, in the second output argument, the gradient value `g`, a vector, at `x`. Note that by checking the value of `nargout` the function can avoid computing `g` when `fun` is called with only one output argument (in the case where the optimization algorithm only needs the value of `f` but not `g`).

```
function [f,g] = myfun(x)
f = ...           % Compute the function value at x
if nargout > 1    % fun called with two output arguments
    g = ...       % Compute the gradient evaluated at x
end
```

The gradient consists of the partial derivatives of `f` at the point `x`. That is, the *i*th component of `g` is the partial derivative of `f` with respect to the *i*th component of `x`. If the Hessian matrix can also be computed *and* the `Hessian` parameter is 'on', i.e., `options = optimset('Hessian','on')`, then the function `fun` must return the Hessian value `H`, a symmetric matrix, at `x` in a third output argument. Note that by checking the value of `nargout` we can avoid computing `H` when `fun` is called with only one or two output arguments (in the case where the optimization algorithm only needs the values of `f` and `g` but not `H`).

```
function [f,g,H] = myfun(x)
f = ...           % Compute the objective function value at x
if nargout > 1    % fun called with two output arguments
    g = ...       % Gradient of the function evaluated at x
    if nargout > 2
        H = ...   % Hessian evaluated at x
    end
end
```

The Hessian matrix is the second partial derivatives matrix of `f` at the point `x`. That is, the (*i,j*)th component of `H` is the second partial derivative of `f` with respect to `xi` and `xj`,

$\frac{\partial^2 f}{\partial x_i \partial x_j}$. The Hessian is by definition a symmetric matrix.

`nonlcon` The function that computes the nonlinear inequality constraints $c(x) \leq 0$ and the nonlinear equality constraints $ceq(x) = 0$. The function `nonlcon` accepts a vector x and returns two vectors c and ceq . The vector c contains the nonlinear inequalities evaluated at x , and ceq contains the nonlinear equalities evaluated at x . The function `nonlcon` can be specified as a function handle.

```
x = fmincon(@myfun,x0,A,b,Aeq,beq,lb,ub,@mycon)
```

where `mycon` is a MATLAB function such as

```
function [c,ceq] = mycon(x)
c = ...      % Compute nonlinear inequalities at x.
ceq = ...    % Compute nonlinear equalities at x.
```

If the gradients of the constraints can also be computed *and* the `GradConstr` parameter is 'on', as set by

```
options = optimset('GradConstr','on')
```

then the function `nonlcon` must also return, in the third and fourth output arguments, `GC`, the gradient of $c(x)$, and `GCEq`, the gradient of $ceq(x)$. Note that by checking the value of `nargout` the function can avoid computing `GC` and `GCEq` when `nonlcon` is called with only two output arguments (in the case where the optimization algorithm only needs the values of c and ceq but not `GC` and `GCEq`).

```
function [c,ceq,GC,GCEq] = mycon(x)
c = ...      % Nonlinear inequalities at x
ceq = ...    % Nonlinear equalities at x
if nargout > 2 % nonlcon called with 4 outputs
    GC = ...  % Gradients of the inequalities
    GCEq = ... % Gradients of the equalities
end
```

If `nonlcon` returns a vector c of m components and x has length n , where n is the length of x_0 , then the gradient `GC` of $c(x)$ is an n -by- m matrix, where $GC(i,j)$ is the partial derivative of $c(j)$ with respect to $x(i)$ (i.e., the j th column of `GC` is the gradient of the j th inequality constraint $c(j)$). Likewise, if `ceq` has p components, the gradient `GCEq` of $ceq(x)$ is an n -by- p matrix, where $GCEq(i,j)$ is the partial derivative of $ceq(j)$ with respect to $x(i)$ (i.e., the j th column of `GCEq` is the gradient of the j th equality constraint $ceq(j)$).

`options` [Options](#) provides the function-specific details for the `options` parameters.

Output Arguments

[Function Arguments](#) contains general descriptions of arguments returned by `fmincon`. This section provides function-specific details for `exitflag`, `lambda`, and `output`:

exitflag	Describes the exit condition:
> 0	The function converged to a solution x .
0	The maximum number of function evaluations or iterations was exceeded.
< 0	The function did not converge to a solution.
lambda	Structure containing the Lagrange multipliers at the solution x (separated by constraint type). The fields of the structure are:
lower	Lower bounds lb
upper	Upper bounds ub
ineqlin	Linear inequalities
eqlin	Linear equalities
ineqnonlin	Nonlinear inequalities
eqnonlin	Nonlinear equalities
output	Structure containing information about the optimization. The fields of the structure are:
iterations	Number of iterations taken.
funcCount	Number of function evaluations.
algorithm	Algorithm used.
cgiterations	Number of PCG iterations (large-scale algorithm only).
stepsize	Final step size taken (medium-scale algorithm only).
firstorderopt	Measure of first-order optimality (large-scale algorithm only). For large-scale bound constrained problems, the first-order optimality is the infinity norm of $v \cdot g$, where v is defined as in Box Constraints , and g is the gradient. For large-scale problems with only linear equalities, the first-order optimality is the infinity norm of the <i>projected</i> gradient (i.e. the gradient projected onto the nullspace of A_{eq}).

Options

Optimization options parameters used by `fmincon`. Some parameters apply to all algorithms, some are only relevant when using the large-scale algorithm, and others are only relevant when using the medium-scale algorithm. You can use [optimset](#) to set or change the values of these fields in the parameters structure, `options`. See [Optimization Parameters](#), for detailed information.

We start by describing the `LargeScale` option since it states a *preference* for which algorithm to use. It is only a preference since certain conditions must be met to use the large-scale algorithm. For `fmincon`, you must provide the gradient (see the description of `fun` above to see how) or else the medium-scale algorithm is used:

`LargeScale` Use large-scale algorithm if possible when set to `'on'`. Use medium-scale algorithm when set to `'off'`.

Medium-Scale and Large-Scale Algorithms. These parameters are used by both the medium-scale and large-scale algorithms:

`Diagnostics` Print diagnostic information about the function to be minimized.

`Display` Level of display. `'off'` displays no output; `'iter'` displays output at each iteration; `'final'` (default) displays just the final output.

`GradObj` Gradient for the objective function defined by user. See the description of [fun](#) above to see how to define the gradient in `fun`. You must provide the gradient to use the large-scale method. It is optional for the medium-scale method.

`MaxFunEvals` Maximum number of function evaluations allowed.

`MaxIter` Maximum number of iterations allowed.

`TolFun` Termination tolerance on the function value.

`TolCon` Termination tolerance on the constraint violation.

`TolX` Termination tolerance on x .

Large-Scale Algorithm Only. These parameters are used only by the large-scale algorithm:

`Hessian` If `'on'`, `fmincon` uses a user-defined Hessian (defined in [fun](#)), or Hessian information (when using `HessMult`), for the objective function. If `'off'`, `fmincon` approximates the Hessian using finite differences.

`HessMult` Function handle for Hessian multiply function. For large-scale structured problems, this function computes the Hessian matrix product $H*Y$ without actually forming H . The function is of the form

$$W = \text{hmfun}(\text{Hinfo}, Y, p1, p2, \dots)$$

where `Hinfo` and the additional parameters `p1, p2, ...` contain the matrices used to compute $H*Y$. The first argument must be the same as the third argument returned by the objective function `fun`.

`[f,g,Hinfo] = fun(x,p1,p2,...)`

The parameters `p1, p2, ...` are the same additional parameters that are passed to `fmincon` (and to `fun`).

`fmincon(fun,...,options,p1,p2,...)`

Y is a matrix that has the same number of rows as there are dimensions in the problem. $W = H*Y$ although H is not formed explicitly. `fmincon` uses `Hinfo` to compute the preconditioner.

Note 'Hessian' must be set to 'on' for `Hinfo` to be passed from `fun` to `hmfun`.

See [Nonlinear Minimization with a Dense but Structured Hessian and Equality Constraints](#) for an example.

<code>HessPattern</code>	Sparsity pattern of the Hessian for finite-differencing. If it is not convenient to compute the sparse Hessian matrix H in <code>fun</code> , the large-scale method in <code>fmincon</code> can approximate H via sparse finite-differences (of the gradient) provided the <i>sparsity structure</i> of H -- i.e., locations of the nonzeros -- is supplied as the value for <code>HessPattern</code> . In the worst case, if the structure is unknown, you can set <code>HessPattern</code> to be a dense matrix and a full finite-difference approximation is computed at each iteration (this is the default). This can be very expensive for large problems so it is usually worth the effort to determine the sparsity structure.
<code>MaxPCGIter</code>	Maximum number of PCG (preconditioned conjugate gradient) iterations (see the <i>Algorithm</i> section below).
<code>PrecondBandWidth</code>	Upper bandwidth of preconditioner for PCG. By default, diagonal preconditioning is used (upper bandwidth of 0). For some problems, increasing the bandwidth reduces the number of PCG iterations.
<code>TolPCG</code>	Termination tolerance on the PCG iteration.
<code>TypicalX</code>	Typical x values.

Medium-Scale Algorithm Only. These parameters are used only by the medium-scale algorithm:

<code>DerivativeCheck</code>	Compare user-supplied derivatives (gradients of the objective and constraints) to finite-differencing derivatives.
<code>DiffMaxChange</code>	Maximum change in variables for finite-difference gradients.
<code>DiffMinChange</code>	Minimum change in variables for finite-difference gradients.

Examples

Find values of x that minimize $f(x) = -x_1x_2x_3$, starting at the point $x = [10; 10; 10]$ and subject to the constraints

$$0 \leq x_1 + 2x_2 + 2x_3 \leq 72$$

First, write an M-file that returns a scalar value f of the function evaluated at x .

```
function f = myfun(x)
f = -x(1) * x(2) * x(3);
```

Then rewrite the constraints as both less than or equal to a constant,

$$\begin{aligned} -x_1 - 2x_2 - 2x_3 &\leq 0 \\ x_1 + 2x_2 + 2x_3 &\leq 72 \end{aligned}$$

Since both constraints are linear, formulate them as the matrix inequality $A \cdot x \leq b$ where

$$A = \begin{bmatrix} -1 & -2 & -2 \\ 1 & 2 & 2 \end{bmatrix} \quad b = \begin{bmatrix} 0 \\ 72 \end{bmatrix}$$

Next, supply a starting point and invoke an optimization routine.

```
x0 = [10; 10; 10]; % Starting guess at the solution
[x,fval] = fmincon(@myfun,x0,A,b)
```

After 66 function evaluations, the solution is

```
x =
    24.0000
    12.0000
    12.0000
```

where the function value is

```
fval =
   -3.4560e+03
```

and linear inequality constraints evaluate to be ≤ 0

```
A*x-b=
   -72
    0
```

Notes

Large-Scale Optimization. To use the large-scale method, the gradient must be provided in `fun` (and the `GradObj` parameter is set to `'on'`). A warning is given if no gradient is provided and the `LargeScale` parameter is not `'off'`. The function `fmincon` permits $g(x)$ to be an approximate gradient but this option is not recommended; the numerical behavior of most optimization codes is considerably more robust when the true gradient is used.

The large-scale method in `fmincon` is most effective when the matrix of second derivatives, i.e., the Hessian matrix $H(x)$, is also computed. However, evaluation of the true Hessian matrix is not required. For example, if you can supply the Hessian sparsity structure (using the `HessPattern` parameter in `options`), then `fmincon` computes a sparse finite-difference approximation to $H(x)$.

If x_0 is not strictly feasible, `fmincon` chooses a new strictly feasible (centered) starting point.

If components of x have no upper (or lower) bounds, then `fmincon` prefers that the corresponding components of `ub` (or `lb`) be set to `Inf` (or `-Inf` for `lb`) as opposed to an arbitrary but very large positive (or negative in the case of lower bounds) number.

Several aspects of linearly constrained minimization should be noted:

- A dense (or fairly dense) column of matrix A_{eq} can result in considerable fill and computational cost.
- `fmincon` removes (numerically) linearly dependent rows in A_{eq} ; however, this process involves repeated matrix factorizations and therefore can be costly if there are many dependencies.
- Each iteration involves a sparse least-squares solve with matrix

$$\overline{A_{eq}} = A_{eq}^T R^{-T}$$

where R is the Cholesky factor of the preconditioner. Therefore, there is a potential conflict between choosing an effective preconditioner and minimizing fill in $\overline{A_{eq}}$.

Medium-Scale Optimization. Better numerical results are likely if you specify equalities explicitly using `Aeq` and `beq`, instead of implicitly using `lb` and `ub`.

If equality constraints are present and dependent equalities are detected and removed in the quadratic subproblem, `'dependent'` is printed under the `Procedures` heading (when you ask for output by setting the `Display` parameter to `'iter'`). The dependent equalities are only removed when the equalities are consistent. If the system of equalities is not consistent, the subproblem is infeasible and `'infeasible'` is printed under the `Procedures` heading.

Algorithm

Large-Scale Optimization. By default `fmincon` will choose the large-scale algorithm *if* the user supplies the gradient in `fun` (and `GradObj` is `'on'` in `options`) *and* if *only* upper and lower bounds exist *or* *only* linear equality constraints exist. This algorithm is a subspace trust region method and is based on the interior-reflective Newton method described in [1], [2]. Each iteration involves the approximate solution of a large linear system using the method of preconditioned conjugate gradients (PCG). See the trust-region and preconditioned conjugate gradient method descriptions in the [Large-Scale Algorithms](#) chapter.

Medium-Scale Optimization. `fmincon` uses a Sequential Quadratic Programming (SQP) method. In this method, a Quadratic Programming (QP) subproblem is solved at each iteration. An estimate of the Hessian of the Lagrangian is updated at each iteration using the BFGS formula (see [fminunc](#), references [\[7\]](#), [\[8\]](#)).

A line search is performed using a merit function similar to that proposed by [\[4\]](#), [\[5\]](#), and [\[6\]](#). The QP subproblem is solved using an active set strategy similar to that described in [\[3\]](#). A full description of this algorithm is found in [Constrained Optimization](#) in "Introduction to Algorithms."

See also [SQP Implementation](#) in "Introduction to Algorithms" for more details on the algorithm used.

Diagnostics

Large-Scale Optimization. The large-scale code does not allow equal upper and lower bounds. For example if `lb(2)==ub(2)`, then `fmincon` gives the error

```
Equal upper and lower bounds not permitted in this large-scale
method.
Use equality constraints and the medium-scale method instead.
```

If you only have equality constraints you can still use the large-scale method. But if you have both equalities and bounds, you must use the medium-scale method.

Limitations

The function to be minimized and the constraints must both be continuous. `fmincon` may only give local solutions.

When the problem is infeasible, `fmincon` attempts to minimize the maximum constraint value.

The objective function and constraint function must be real-valued, that is they cannot return complex values.

Large-Scale Optimization. To use the large-scale algorithm, the user must supply the gradient in `fun` (and `GradObj` must be set 'on' in `options`), and only upper and lower bounds constraints may be specified, *or only* linear equality constraints must exist and `Aeq` cannot have more rows than columns. `Aeq` is typically sparse. See [Table 2-4, Large-Scale Problem Coverage and Requirements](#), for more information on what problem formulations are covered and what information must be provided.

Currently, if the analytical gradient is provided in `fun`, the `options` parameter `DerivativeCheck` cannot be used with the large-scale method to compare the analytic gradient to the finite-difference gradient. Instead, use the medium-scale method to check the derivative with `options` parameter `MaxIter` set to 0 iterations. Then run the problem with the large-scale method.

See Also

[@\(function_handle\)](#), [fminbnd](#), [fminsearch](#), [fminunc](#), [optimset](#)

References

[1] Coleman, T.F. and Y. Li, "An Interior, Trust Region Approach for Nonlinear Minimization Subject to Bounds," *SIAM Journal on Optimization*, Vol. 6, pp. 418-445, 1996.

- [2] Coleman, T.F. and Y. Li, "On the Convergence of Reflective Newton Methods for Large-Scale Nonlinear Minimization Subject to Bounds," *Mathematical Programming*, Vol. 67, Number 2, pp. 189-224, 1994.
- [3] Gill, P.E., W. Murray, and M.H. Wright, *Practical Optimization*, Academic Press, London, 1981.
- [4] Han, S.P., "A Globally Convergent Method for Nonlinear Programming," *Journal of Optimization Theory and Applications*, Vol. 22, p. 297, 1977.
- [5] Powell, M.J.D., "A Fast Algorithm for Nonlinear Constrained Optimization Calculations," *Numerical Analysis*, ed. G.A. Watson, *Lecture Notes in Mathematics*, Springer Verlag, Vol. 630, 1978.
- [6] Powell, M.J.D., "The Convergence of Variable Metric Methods For Nonlinearly Constrained Optimization Calculations," *Nonlinear Programming 3* (O.L. Mangasarian, R.R. Meyer, and S.M. Robinson, eds.) Academic Press, 1978.

 fminbnd

fminimax 